

Formalizing Coq Modules in the MetaCoq project

XFC4101 CA Report

Yee-Jian Tan

October 31, 2022

Contents

| | |
|---|------------|
| Contents | iii |
| 1. Introduction | 1 |
| 2. Types and Proof Assistants | 2 |
| 2.1. What are proof assistants? | 2 |
| 2.2. Type Theory | 2 |
| 2.3. The Curry-Howard Correspondence | 2 |
| 2.3.1. Intuitionistic Propositional Logic | 3 |
| 2.3.2. Simply Typed Lambda Calculus | 4 |
| 2.3.3. The Correspondence | 4 |
| 3. Introduction to the MetaCoq Project | 5 |
| 3.1. The MetaCoq Project | 5 |
| 3.2. Structure of the MetaCoq Project | 5 |
| 3.2.1. TemplateCoq | 5 |
| 3.2.2. PCUIC | 6 |
| 3.2.3. Safe Checker, Erasure and Beyond | 6 |
| 4. The Coq Proof Assistant | 7 |
| 4.1. An example of a Coq program | 7 |
| 4.2. Coq Modules | 7 |
| 4.3. A specification of Coq Modules | 9 |
| 4.3.1. Conversion of Coq Terms | 9 |
| 4.3.2. Modules as second-class objects | 9 |
| 4.3.3. Related Works | 10 |
| 4.4. Semantic of Modules | 10 |
| 4.4.1. Plain Modules | 10 |
| 4.4.2. Aliased Modules | 11 |
| 5. Project Plan for The Next Semester | 12 |
| 5.1. Implement Plain Coq Modules | 12 |
| 5.2. Study Type Theoretic Results | 12 |
| APPENDIX | 14 |
| A. Specification of Coq Modules | 16 |
| A.1. Typing Rules | 16 |
| B. The Untyped Lambda Calculus | 17 |
| B.1. The Church-Rosser Confluence | 17 |
| B.2. Cut elimination | 17 |
| Bibliography | 18 |

The Coq Proof Assistant is one of the prominent proof assistants, with applications from the Homotopy Type Theory (HoTT)'s univalent foundation of mathematics, to the Iris proof mode for Concurrent Separation Logic, or the recent recipient of the 2021 ACM Software System Award — the CompCert compiler. However, even though the theory behind Coq, the Calculus of Inductive Constructions is known to be consistent and strongly normalizing (hence proof-writing is decidable), but the OCaml implementation of Coq is known to have an average of one critical bug per year which allows one to prove False statements in Coq.

The MetaCoq project is therefore started by the Galinette Team in INRIA, to "formalize Coq in Coq" and acts as a platform to interact with Coq's terms directly, in a verified manner. This also reduces the trust in the implementation of Coq to the correctness of the theories underlying Coq. In 2020, the core language of Coq, minus a few features are already successfully verified in Coq. However, there are still a few missing pieces not yet verified, among them the Module system of Coq.

The Module system, although not part of the core calculus of Coq, is an important feature for Coq developers to develop modularly, providing massive abstraction and a suitable interface for reusing definitions and theorems.

In this project, I aim to formalize the Module system of Coq using the MetaCoq project framework, by first understanding the implementation of Modules, and then providing a specification of the implementation of Coq modules, finally writing proofs to show the correctness of the current implementation. In particular, I will focus on the formalization of non-parametrized, plain modules.

In order to tackle this project, it is important to understand the theory behind the implementation of Coq; more explicitly, how Type Theory helps in theorem proving. This report is therefore structured from the bottom up, from some results of Type Theory which I study under the supervision of Professor Yang Yue; to the formalization of Modules jointly supervised by Professor Nicolas Tabareau and Professor Martin Henz.

The first section will give an overview of Type Theory and Theorem Proving. The second section will cover some important results from Type Theory that I studied. Following that, I will give an introduction to the language of Coq and the core problem of formalizing Coq Modules, before describing some related works and finally a specification for Coq Modules which I will implement.

2. Types and Proof Assistants

2.1. What are proof assistants?

Proof assistants are a special kind of computer program whose job is to verify the correctness of a mathematical proof, that is, verify if a conclusion can be made by applying a fixed set of logical deduction steps on a set of fixed assumptions. Modern implementations of proof assistants usually utilizes theories such as Higher Order Logic (eg. Isabelle, HOL Light), or some form of Type Theory.

In order to have a computer program that can do such powerful logical reasoning to contain sophisticated mathematical results, the first question that arises is how to represent a theorem and what qualifies as a valid proof of a theorem. Fortunately, Type Theory gives a simple and direct relationship between Logic and Types.

2.2. Type Theory

Types were first invented by Bertrand Russell as a way to provide higher-order structures than sets to overcome his famous Russell's paradox. However, it was Alonzo's Church's Lambda Calculus which when used with types (Simply Typed Lambda Calculus) exhibited many desirable properties as not just a consistent formal system, also as a rewriting system that exhibits strong normalization. The arguably most important characteristic of Type Theories to qualify as a useful theory for the implementation of proof assistants, lies in the Curry-Howard Correspondence.

2.3. The Curry-Howard Correspondence

The Curry-Howard Correspondence establishes an one-to-one correspondence between two distinct disciplines — Logic and Types. It asserts that there is a one-to-one correspondences between Types and Propositions, and between Terms and Proofs. More precisely, it states that every proposition in intuitionistic propositional logic (constructive 0-th order logic) can be expressed as a Type in Simply Typed Lambda Calculus, and correspondingly, a term of a specific type is viewed as a proof for the proposition represented by the type.

| Intuitionistic Propositional Logic | Type Theory |
|------------------------------------|-------------|
| Proposition | Type |
| Proof | Term |

What are those exactly?

2.3.1. Intuitionistic Propositional Logic

Propositional logic, also known as *0-th order logic* is a logical system where propositions (sentences) are constructed with variables and logical connectors, such as \rightarrow (implication), \vee (disjunction), \neg (negation), \wedge (conjunction) and more. For example, propositions in Propositional logic looks like:

1. $p \vee q$
2. $p \wedge \neg p$
3. $(p \rightarrow q) \rightarrow (\neg q \rightarrow \neg p)$

One can assign a valuation to each variable, either 1 or 0, which represents "True" or "False" informally. If there is some valuation that result in the proposition result in 1, then we say the proposition is *satisfiable*; if there is none, i.e. all valuations will result in the proposition being 0, we say the proposition is a *contradiction*; on the other hand, if any valuation will result in 1, the proposition is a *tautology*.

In particular, in the above example, the propositions are satisfiable, a contradiction and a tautology in that order.

Logical systems can be either intuitionistic or classical, based on one fact: whether they accept the *rule of the excluded middle*. Precisely, it says

Axiom 2.3.1 (Excluded Middle) For any proposition p , either p or $\neg p$ holds.

Much of modern mathematics is built on classical logic, which asserts the axiom of excluded middle. This leads to the existence of multitude of non-constructive proofs, which shows existence without giving a witness. A classical non-constructive proof using the law of excluded middle is as follows:

Theorem 2.3.1 *There exists irrational numbers a, b where a^b is rational.*

Proof. By the law of excluded middle, $\sqrt{2}^{\sqrt{2}}$ is either rational or irrational.

Suppose it is rational, then take $a = b = \sqrt{2}$. We have found such a, b .

Suppose it is irrational, then take $a = \sqrt{2}^{\sqrt{2}}$, $b = \sqrt{2}$. Then

$$a^b = \left(\sqrt{2}^{\sqrt{2}} \right)^{\sqrt{2}} = \sqrt{2}^{\sqrt{2} \times \sqrt{2}} = 2$$

is rational. □

Even with the proof, we still don't have a concrete pair of irrational numbers (a, b) fulfilling the property, but the conclusion is declared to be logically sound.

Intuitionistic logic, on the other hand, rejects this rule, leading to a logical system where to prove something exists, it requires an algorithm for the explicit construction of such object.

2.3.2. Simply Typed Lambda Calculus

The Simply Typed Lambda Calculus (STLC) is a simple construction:

2.3.3. The Correspondence

This correspondence is two-way: in particular, we can encode a proposition as a type, and to write a proof is equivalent to finding a term that inhabit that particular type. This reduced the problem of proof-checking to the problem of finding the type of a term, or under the Curry-Howard correspondence, whether a proof truly proves a given proposition.

This fundamental correspondence is elegant, but only limited to propositional logic. More complex type systems, such as the Calculus of Constructions by Thierry Coquand, extend the Curry-Howard correspondence to higher-order logic, finally allowing the encoding of most mathematical truth and specifying proofs for theories in those systems.

Introduction to the MetaCoq Project

3.

3.1. The MetaCoq Project

MetaCoq is a project to formalize the core calculus, PCUIC, in Coq, and become a platform to write tools that can manipulate Coq terms. The effort was complete for a large part of the core language of Coq, with a few missing pieces:

- ▶ Eta
- ▶ Template Polymorphism
- ▶ SProps
- ▶ Modules

I will be tackling the last.

3.2. Structure of the MetaCoq Project

The MetaCoq project is an ambitious project aiming to provide a verified implementation of Coq, and for its size, it is reasonably split into a few main components. From the layer closest to the Coq language to the machine code, we have: TemplateCoq (Section 3.2.1), PCUIC (Section 3.2.2), followed by Safe Checker, Erasure and beyond (Section 3.2.3).

Let us remind ourselves of the task of MetaCoq: we would like to see that the OCaml representation of Coq is indeed correct and preserves the desired properties of the underlying theory. Since Coq has added many bells and whistles for its users, the terms of Coq definitely is much more complex than its underlying, platonic type theoretical form. Therefore, MetaCoq has several stages for a term to go through, stripping down to the bare minimum through the following few stages.

3.2.1. TemplateCoq

TemplateCoq is a quoting library for Coq: a Coq program that takes a Coq term, and constructs an inductive data type that correspond to its kernel representation in the OCaml implementation. This is the first layer of the stripping of a Coq term, where the structures are preserved properly.

This allows one to turn a Coq program into a Coq internal representation along with its associated environment structures, such as the definitions and declarations in the environment.

3.2.2. PCUIC

PCUIC is the Polymorphic Cumulative Calculus of Inductive Constructions. It is a "cleaned up version of the term language of Coq and its associated type system, shown equivalent to the one in Coq." (From MetqCoq website). In other words, it is a type theory that is as powerful as Coq can express, having the good properties such as weakening, confluence, principality (that every term has a principal type) etc.

A term generated in TemplateCoq can be converted into PCUIC term via a verified process. The theory of PCUIC is then proven to have all the nice properties in Coq.

3.2.3. Safe Checker, Erasure and Beyond

The core semantic operation of type theories are the reductions. The safechecker is a verified "reduction machine, conversion checker and type checker" for PCUIC terms. So far, we have the tools to start with a Coq term, first quoting into TemplateCoq, then converted into a PCUIC term, and eventually has its type checked in the Safe Checker via a fully verified process. As far as correctness is concerned, this has already formed a verified end-to-end process of Coq's correctness.

The MetaCoq has further provided a verified Type and Proof erasure process from PCUIC to untyped Lambda Calculus. The equivalence of this erased language is can be evaluated in *C-light* semantics, the subset of C accepted by the CompCert verified compiler, which completes a maximally safe evaluation toolchain for the language of Coq [1].

[1]: Sozeau et al. (2019), *Coq Coq Correct! Verification of Type Checking and Erasure for Coq, in Coq*

The Coq Proof Assistant

4.

The Coq Proof Assistant (or Coq for short) is a proof assistant based on the Calculus of Inductive Constructions, which added co-inductive types into the Calculus of Constructions.

4.1. An example of a Coq program

This is a simple example featuring the definition of a type “nat”, a definition of a term of type “nat”, namely “zero”, and a function “plus” which takes two “nat”s and returns its sum.

```
Inductive nat: Set :=  
  | 0  
  | S : nat -> nat.  
  
Definition zero : nat := 0.  
  
Fixpoint plus (n m: nat) : nat :=  
  match n with  
  | S n' => S (plus n' m)  
  | 0 => m  
  end.  
  
Proposition zero_is_left_additive_identity : Prop:  
  forall n: nat, plus zero n = n.  
Proof.  
  intro n.  
  simpl.  
  reflexivity.  
Qed.
```

Listing 1: A simple Coq program.

Finally, as explained before, a proposition is a type; and here a named proposition is postulated. A simple proof is also given, which constructs a term inhabiting that type, hence proving the correctness of this proposition.

4.2. Coq Modules

Modules in Coq not only allows the reuse of code, it also provides parametrized theories or data structures in the form of functors (or parametrized modules).

In the following example, we show how one can package definitions into named modules for reuse; we also show how we can create parametrized generic data structures.

In the final line, a new Magma was created by the functor “DoubleMagma”.

Listing 2: An example of Modules.

```

Inductive nat :=
| 0
| S : nat -> nat.

Fixpoint plus (n m: nat) :=
  match n with
  | S n' => S (plus n' m)
  | 0 => m
  end.

(* A magma is a set with a binary (closed) operation. *)
Module Type Magma.
  Parameter T: Set.
  Parameter op: T -> T -> T.
End Magma.

(* Natural numbers with plus form a magma. *)
Module Nat: Magma.
  Definition T := nat.
  Definition op := plus.
End Nat.

(* A functor transforming a magma into another magma. *)
Module DoubleMagma (M: Magma): Magma.
  Definition T := M.T.
  Definition op x y := M.op (M.op x y) (M.op x y).
End DoubleMagma.

Module NatWithDoublePlus := DoubleMagma Nat.

```

4.3. A specification of Coq Modules

Module systems are a feature of the ML family of languages; it allows for massive abstraction and the reuse of code. In particular, Coq also has a module system that is influenced by ML modules, first implemented by Jacek Chrząszcz in 2003, then modified by Elie Soubrian in 2010.

There are a few keywords when it comes to Coq Modules:

- ▶ A **structure** is a anonymous collection of definitions, and is the underlying construct of modules. They contain **structure elements**, which can be a
 - constant definition


```
Definition a: bool := true.
```
 - assumption


```
Axiom inconsistent: forall p: Prop, p.
```
 - **module, module type, functors** recursively.
- ▶ A **module** is a structure given a name. It can be defined explicitly to be of a certain **module type**, which a named structure with possibly empty definitions.
- ▶ A **module alias** is the association of a short name to an existing module.
- ▶ A **functor** is a module defined with a parameter with a binder and a required type for the module supplied as an argument.

For a more precise definition of modules and related structures, please refer to [Coq: Modules](#).

4.3.1. Conversion of Coq Terms

To understand Coq Modules, we need to first understand the basic structure of Coq. The core object in the language of Coq are terms. Terms of a type correspond to a proof for a theorem as in the Curry-Howard correspondence. The syntax and semantics of Coq terms are as explained by the syntax,¹ conversion (including reduction and expansion)² and typing³ respectively. The evaluation of Coq terms are done under a Global Environment Σ containing definitions, and a local context Γ containing assumptions. Evaluation in Coq is known as conversion, the reflexive, transitive closure of the various reduction rules that is defined, including the famous β -reduction (function application).

- 1: [Coq: Essential Vocabulary](#)
- 2: [Coq: Conversion](#)
- 3: [Coq: Typing](#)

The conversion relation is then defined with these parameters, from which we can conclude nice properties on conversion, such as strong normalization, confluence and decidability.

4.3.2. Modules as second-class objects

However, Coq modules are not first class objects of the language and do not participate in conversion themselves; i.e. there is no notion to "reduce" a module. Plain modules in Coq can be treated as a named container of constant and inductive definitions, including possibly nested modules; namespaced by a dot-separated string called a "path". This abstraction

allows users to reuse definitions, essentially importing another "global environment" into the current one.

To further expand this possibility, module functors exist to be interfaces which users can provide definitions for, by supplying a module definition. Functors are therefore opaque second-class objects which is only useful when a module is generated.

In this chapter, we formalize the syntax and semantics of the current implementation of Coq and implement in the level of TemplateCoq. Section 4.4.1 describes the implementation of plain modules without functors.

4.3.3. Related Works

[2]: Chrzyszcz (2003), *Implementing Modules in the Coq System*

[3]: Soubiran (2010), *Développement modulaire de théories et gestion de l'espace de nom pour l'assistant de preuve Coq.*

Jacek Chrzyszcz's article in TPHOLs 2003 [2] explains the motivation and choices for the implementation of modules as a second class object and its interaction with terms as described above. Elie Soubiran's PhD Thesis on Coq Modules [3] describes an extension of the Module system which only some features are implemented. The [ModuleSystem Wiki Page](#) on Coq's official Github repository contains valuable information on the usage and design decision of the Module system of Coq, including a list of open issues with Modules. In general, the issues do not compromise the correctness of the type system implemented by Coq; instead, they should be viewed as possible areas of improvements for Coq users.

[4]: Dreyer et al. (2005), *Understanding and Evolving the ML Module System*

[5]: Rossberg et al. (2010), *F-Ing Modules*

[6]: Tan et al. (2015), *A Verified Type System for CakeML*

[7]: MacQueen (1984), *Modules for standard ML*

Other slightly useful references include papers that explain modules in the ML family of languages, specifically OCaml and to a certain extent, Standard ML. On this note, Derek Dreyer wrote his PhD thesis [4] on understanding and extending ML modules, and subsequently a on implementing ML modules in its most desirable form, applicative and first-order as a subset of a small type system, F_ω [5]. Other notable implementations include that of CakeML [6], a verified ML language and Standard ML modules by MacQueen [7]. However, since the type system of Coq is much stronger and sophisticated compared to ML languages, the implementations also vary wildly and one can only refer to them for inspirations.

4.4. Semantic of Modules

4.4.1. Plain Modules

Plain modules are just saving declarations and definitions in a structure, in the global context. Its contents are not modified during conversion/reduction.

In Coq, modules are second-class objects; in other words, a module is not a term. Its definition is stored and referred to by a canonical kname. We say the implementation of such a module is correct if the metatheory of the original system are unchanged and remains correct; that is the proofs go through when terms can be defined within modules. Since the MetaCoq project has proven various nice properties about conversion in Coq, our project on plain modules is two-fold:

1. Ensure the correctness of the static semantics of Coq Modules (well-typedness).
2. Define the behaviour of access of definitions within Modules.

Once these two are done, we can be sure that a Coq program with Modules has all its terms well-defined (by (1)) and enjoys the nice properties of conversion, since the additional terms defined in Modules fulfill (2). This follows as our definition of Modules on the TemplateCoq level, is eventually elaborated down into the PCUIC calculus the idea of modules and aliasing do not exist anymore, they are flattened into the corresponding global environment.

Concretely, if a module as below is defined while the Global Environment, which stores definitions is denoted as Σ :

```
Module M.
  Definition a: nat := 0.
End M.
```

Then the environment will have an new declaration added:

$$\Sigma := \Sigma :: \text{ModuleDeclaration}(M, [\text{ConstantDeclaration}(M.a, \text{nat}, 0)])$$

So when $M.a$ is called, it refers to the definition in the

4.4.2. Aliased Modules

Aliased modules are just a renaming of existing modules, which can be seen as syntactic sugar for modules. Therefore, the correctness depends only on implementing this internal referencing correctly. Suppose we have

```
Module N := M.
```

Aliasing N to M , then any access path $N.X$ should be resolved similarly to $M.X$ (note that since M is possibly an alias as well, we do not require $N.X$ to resolve to $M.X$). In the OCaml implementation, this is implemented by an internal kernel name, conceptually similar to a lookup table; formally, the kernel name is a pair of a *name* and a *label*, with label being the internal reference used by both N and M .

5. Project Plan for The Next Semester

Since this is a joint thesis with the Mathematics department, also due to personal interest, for the remaining semester for this thesis, I am planning to:

5.1. Implement Plain Coq Modules

There are a few steps to this implementation that I can foresee:

1. Study the structure and conventions of the MetaCoq project. Since this is a huge project, many common theorems have been abstracted into multiple functors which can generate a theorem given a corresponding object, such as
2. Add the data structure of Plain Modules into the global environment which is quoted by TemplateCoq. This requires the modification of the current data structure containing declarations (minus modules).
3. Ensure the environment remains well-defined with the current typing rules and correctness assertions. This requires the modifications to the existing proofs on the environment (well-formedness, well-typedness).
4. Extend the TemplateCoq plugin (in OCaml) which transforms a Coq term into a TemplateCoq representation, to translate a Coq module into the internal representation of Modules I implemented above. This requires OCaml knowledge specifically in the area of Coq plugins.
5. Now that syntactical transformation from Coq to TemplateCoq is complete, I need to define how these modules that live in the Global Environment will be used. In particular, I need to implement the "canonical name" model as implemented in the Coq kernel.
6. Finally, the translation from TemplateCoq to PCUIC. In PCUIC, modules will cease to exist and references to definitions in modules will be treated as direct references to plain definitions outside modules. This is done by referring module paths to their canonical paths.
7. The last step in this project is to verify that the translation from TemplateCoq to PCUIC (and vice versa) is correct; that is, a PCUIC term is well-formed and well-typed iff its translation is well-formed and well-typed in TemplateCoq.

5.2. Study Type Theoretic Results

The modern field of Type Theory and Proof Assistants grow very quickly, with many new fields of logic (modal logic, linear logic, separation logic) etc and their corresponding Type Theory being studied and made into proof assistants to verify results in these fields. Since there were not many relevant courses in my undergraduate curriculum, I decide to continue

to study some classical results alongside with the formalization project above.

The direction which I have chosen is towards a result important for the implementation of proof assistants: the decidability of conversion (correspondingly proof-checking) in a simpler type system, such as the Martin-Lof Type Theory. So far, I have covered results from Simply Typed Lambda Calculus including confluence and weak normalization, and its counterparts in proof theory, namely sequent calculus, natural deduction as a normal form, and cut-elimination.

APPENDIX

A. Specification of Coq Modules

A.1. Typing Rules

I will be typesetting the rules here, taken from [Coq: Typing Modules](#).

The Untyped Lambda Calculus

B.

Here, I present some properties of the untyped lambda calculus I studied. We follow the text "Proof and Types" by Girard. Here, I present some fundamental results.

B.1. The Church-Rosser Confluence

We study the paper by Takahashi's successors on a proof for Church-Rosser theorem. It states that:

Theorem B.1.1 (Church-Rosser) *If $M \rightarrow_{\beta} M_1$, $M \rightarrow_{\beta} M_2$, then there exists a term N such that $M_1 \rightarrow_{\beta} N$ and $M_2 \rightarrow_{\beta} N$.*

This property asserts the uniqueness of a normal form of a rewriting system, if there exists one. This is especially important for systems that are known to normalize; in particular, the Simply Typed Lambda Calculus, the Martin Lof Type Theory, and Calculus of Constructions are all known to be strongly normalizing. Together with the decidability of conversion, we have a viable theory for the implementation of proof assistants.

The crucial step of the proof is to define a notion of parallel reduction, and the key lemma is

Lemma B.1.2 *If $M \rightarrow_{\beta n} N$, then $M \rightarrow_{\beta} M^{n*}$.*

This provides a candidate of confluence just based on the original term M , instead of the intermediate steps during $M \rightarrow_{\beta} M_1, M_2$ or even M_1, M_2 itself.

(Hopefully I have a Coq proof.)

B.2. Cut elimination

Sequent Calculus is a logical system capable of doing deduction. One of the axioms, the Cut rule, is actually eliminable. This can be proven by induction.

(I will insert the proof here).

Bibliography

Here are the references in citation order.

- [1] Matthieu Sozeau et al. 'Coq Coq Correct! Verification of Type Checking and Erasure for Coq, in Coq'. In: *Proc. ACM Program. Lang.* 4.POPL (2019). doi: [10.1145/3371076](https://doi.org/10.1145/3371076) (cited on page 6).
- [2] Jacek Chrząszcz. 'Implementing Modules in the Coq System'. In: (Dec. 2003) (cited on page 10).
- [3] Elie Soubiran. 'Développement modulaire de théories et gestion de l'espace de nom pour l'assistant de preuve Coq.' Theses. Ecole Polytechnique X, Sept. 2010 (cited on page 10).
- [4] Derek Dreyer, Robert Harper, and Karl Crary. 'Understanding and Evolving the ML Module System'. AAI3166274. PhD thesis. USA, 2005 (cited on page 10).
- [5] Andreas Rossberg, Claudio V. Russo, and Derek Dreyer. 'F-Ing Modules'. In: *Proceedings of the 5th ACM SIGPLAN Workshop on Types in Language Design and Implementation*. TLDI '10. Madrid, Spain: Association for Computing Machinery, 2010, pp. 89–102. doi: [10.1145/1708016.1708028](https://doi.org/10.1145/1708016.1708028) (cited on page 10).
- [6] Yong Kiam Tan, Scott Owens, and Ramana Kumar. 'A Verified Type System for CakeML'. In: *Proceedings of the 27th Symposium on the Implementation and Application of Functional Programming Languages*. IFL '15. Koblenz, Germany: Association for Computing Machinery, 2015. doi: [10.1145/2897336.2897344](https://doi.org/10.1145/2897336.2897344) (cited on page 10).
- [7] David MacQueen. 'Modules for standard ML'. In: *Proceedings of the 1984 ACM Symposium on LISP and Functional Programming*. 1984, pp. 198–207 (cited on page 10).